

Data X

Introduction To Convolution Neural Networks

Ikhlaq Sidhu, Chief Scientist, SCET/IEOR, UC Berkeley
Sana Iqbal, GSI, Data-X, UC Berkeley

Review: A simple Feed Forward Neural Net

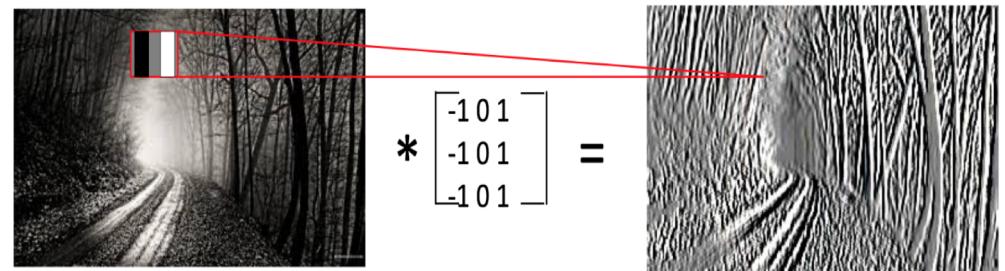
1. Accepts the Input.
2. Creates different combinations of **weighted input**
3. Induces **Non-linearity** in the input using Activations.
4. Repeats 1,2,3 in **hidden layers**.
5. Gives the Output.



Review: Convolution functions in Image Processing

- **Convolution functions** applied between an image matrix and a filter matrix gives us information about the spatial distribution of the image features.
- Using **different types of filters** affects the information extracted from an image- blur, edges, texture.

Convolutional of Two Signals



$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1, n_2] \cdot g[x-n_1, y-n_2]$$

elementwise multiplication and sum



Before We Go On

X is an image



h is filter (or kernel)

$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \rightarrow$$

What is $X * h$?

Data X

A horizontal bar composed of binary digits (0s and 1s) representing the pixel data of the image X. The data is highly compressed and noisy, showing the underlying binary representation of the bread rolls image.

Before We Go On

X is an image



h is filter (or kernel)

$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \rightarrow$$

What is $X * h$?



Answer:

1. It is another image
2. Depending on the kernel, it highlights specific features



Many Operations with Convolution

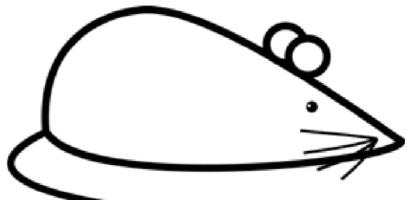
These filters are more typical for image processing,
but not feature extraction.

In CNNs, we are looking to extract features

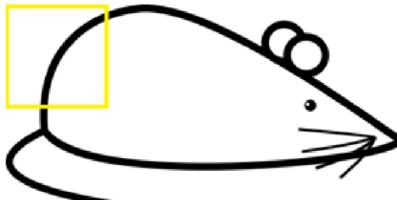
Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	



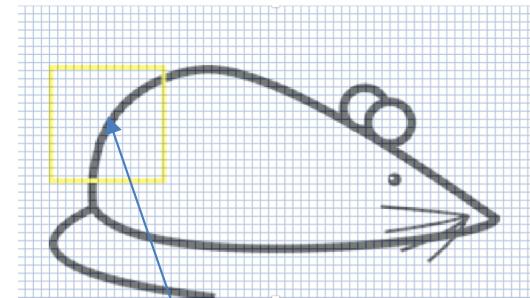
Feature Extraction



Original image



Visualization of the filter on the image



Visualization of the receptive field

0	0	0	0	0	0	30	
0	0	0	0	50	50	50	
0	0	0	20	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	

Pixel representation of the receptive field

*

0	0	0	0	0	0	30	0
0	0	0	0	30	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Pixel representation of filter

$$\text{Multiplication and Summation} = (50*30) + (50*30) + (50*30) + (20*30) + (50*30) = 6600 \text{ (A large number!)}$$

Value of convolution at this spot is large because the sum of products

Think of this like a “moving dot product” of the kernel across the original image

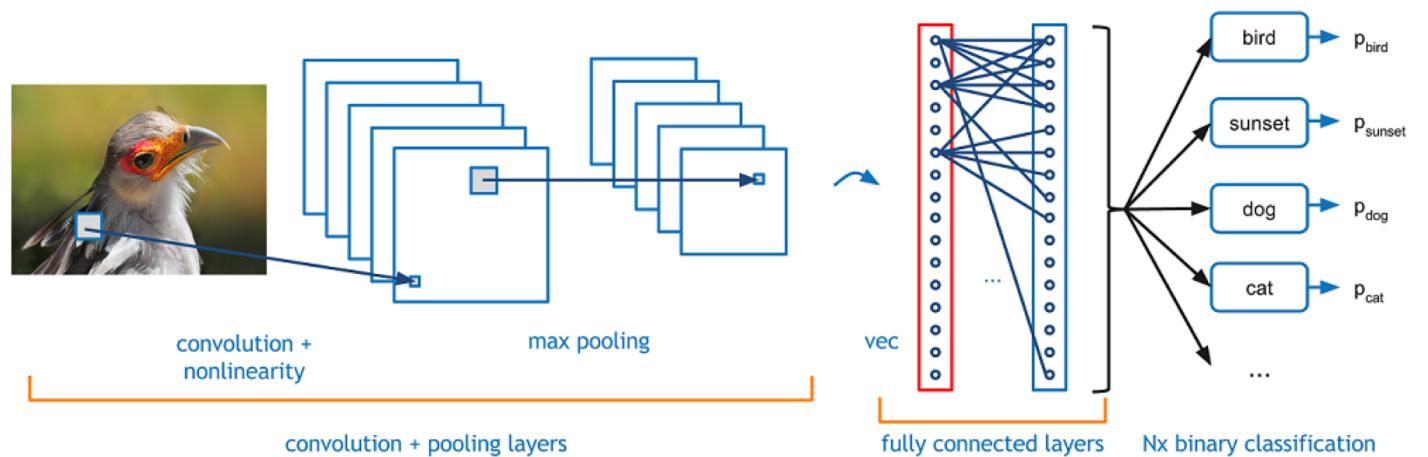
<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner%27s-Guide-To-Understanding-Convolutional-Neural-Networks/>

Data X

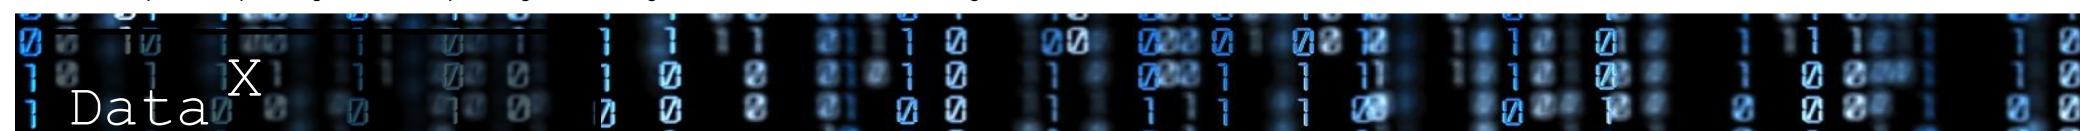
Convolution Neural Networks

So there are four main **operations** in a basic CNN:

1. Convolution
2. Non Linearity (ReLU, tanh, sigmoid)
3. Pooling or Down-Sampling of features
4. Prediction



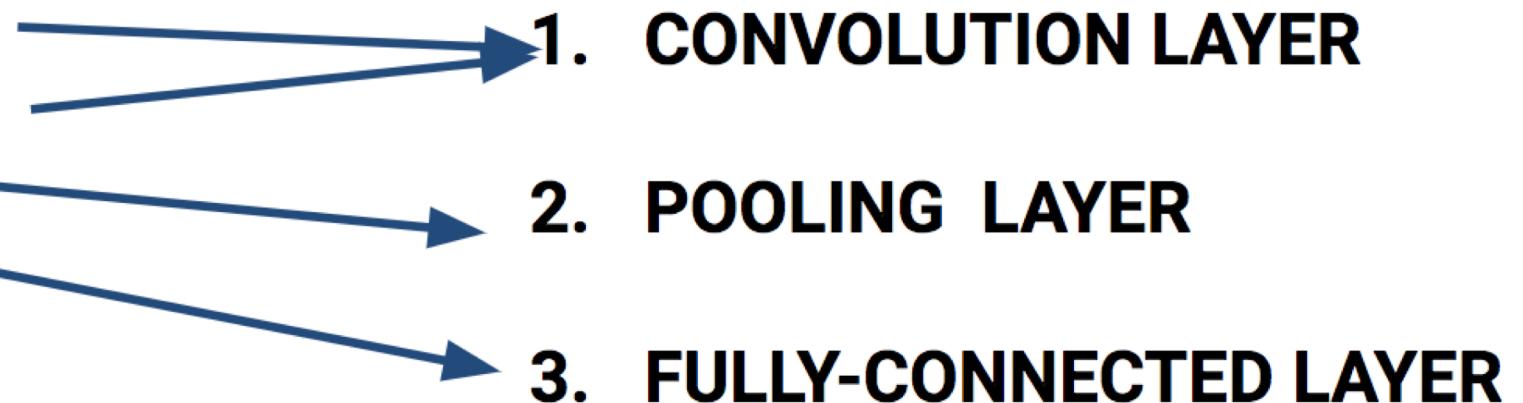
<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner%27s-Guide-To-Understanding-Convolutional-Neural-Networks/>



Convolution Neural Networks

Three different **types of layers** in CNN:

1. Convolution
2. Non Linearity
3. Pooling
4. Prediction



1. CONVOLUTION LAYER

- A. This layer has the filters that we want to apply on our image
- B. Each filter also called kernel has same depth as the channel of the image

For an image with **one channel** the below filter works as:

Sizes:

Image of $h \times w$

Kernel of $f \times g$

New image: $(h-f+1) \times (w-g+1)$

1	0	1
0	1	0
1	0	1

1 $\times 1$	1 $\times 0$	1 $\times 1$	0	0
0 $\times 0$	1 $\times 1$	1 $\times 0$	1	0
0 $\times 1$	0 $\times 0$	1 $\times 1$	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

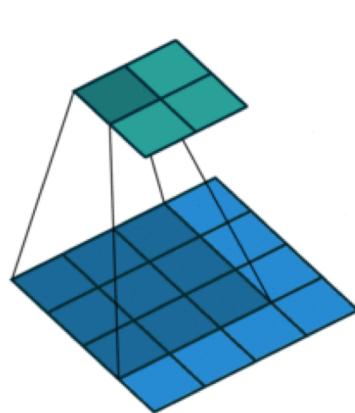
Convolved Feature



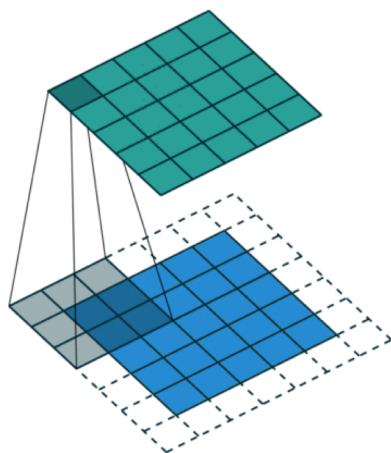
1. CONVOLUTION LAYER

Hyperparameters that can be tuned to change complexity and size of extracted features

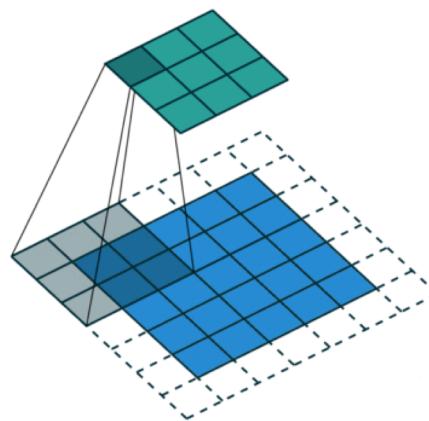
Filter Size, **Stride** (step size of the filter), **Zero-padding** (to maintain dimensions)



- Stride 1
- No Zero padding
- Kernel size 3x3
- Input size 4x4
- Output size 2x2



- Stride 1
- Zero padding 1
- Kernel size 3x3
- Input size 5x5
- Output size 5x5



- Stride 2
- Zero padding 1
- Kernel size 3x3
- Input size 5x5
- Output size 3x3

$$n_{out} = \left\lceil \frac{n_{in} + 2p - k}{s} \right\rceil + 1$$

n_{in} : number of input features

n_{out} : number of output features

k : convolution kernel size

p : convolution padding size

s : convolution stride size

$n_{in} = n_{out}$ when

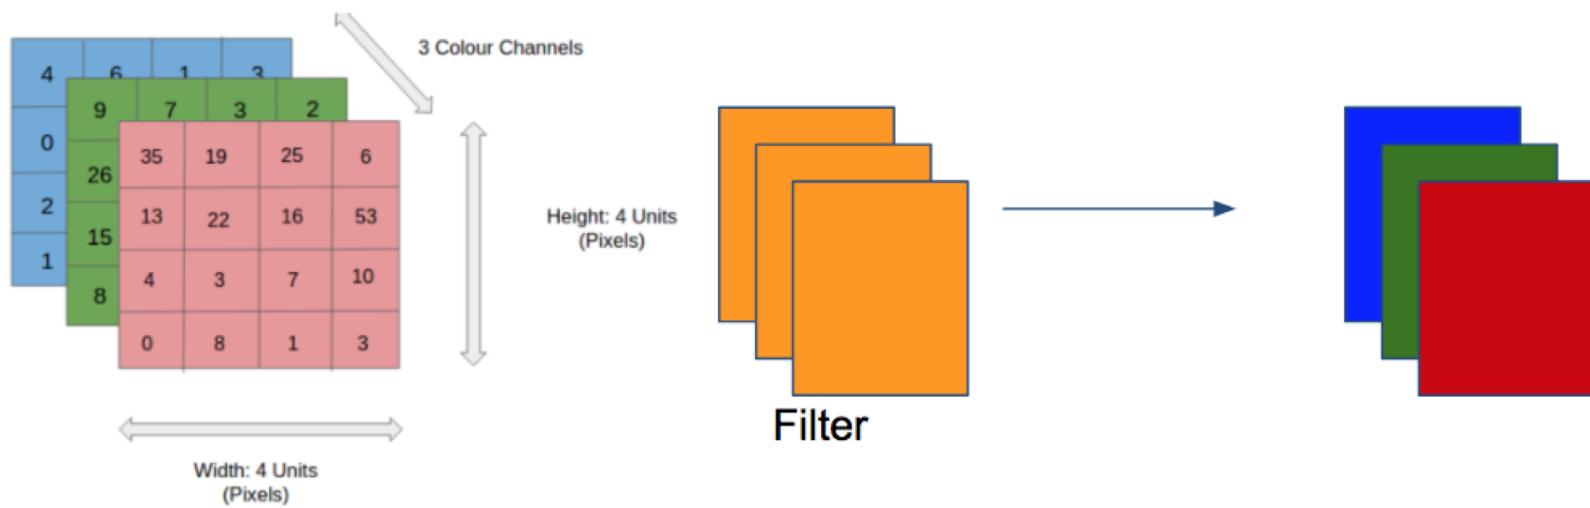
$$\text{Zero Padding} = \frac{(K - 1)}{2}$$

(K = kernel size)

Image Source: https://github.com/vdumoulin/conv_arithmetic

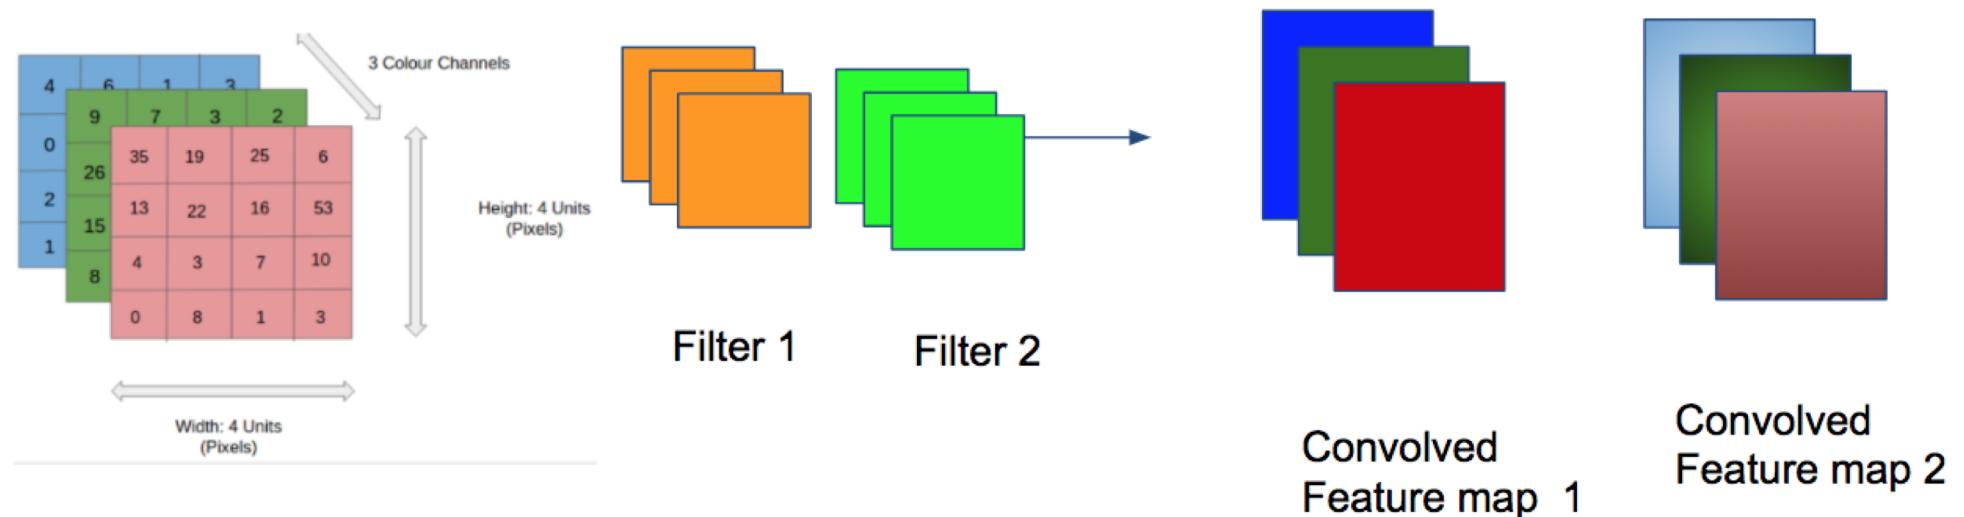
1. CONVOLUTION LAYER

For an image with **3 channels** there would be **1 filter** for each channel

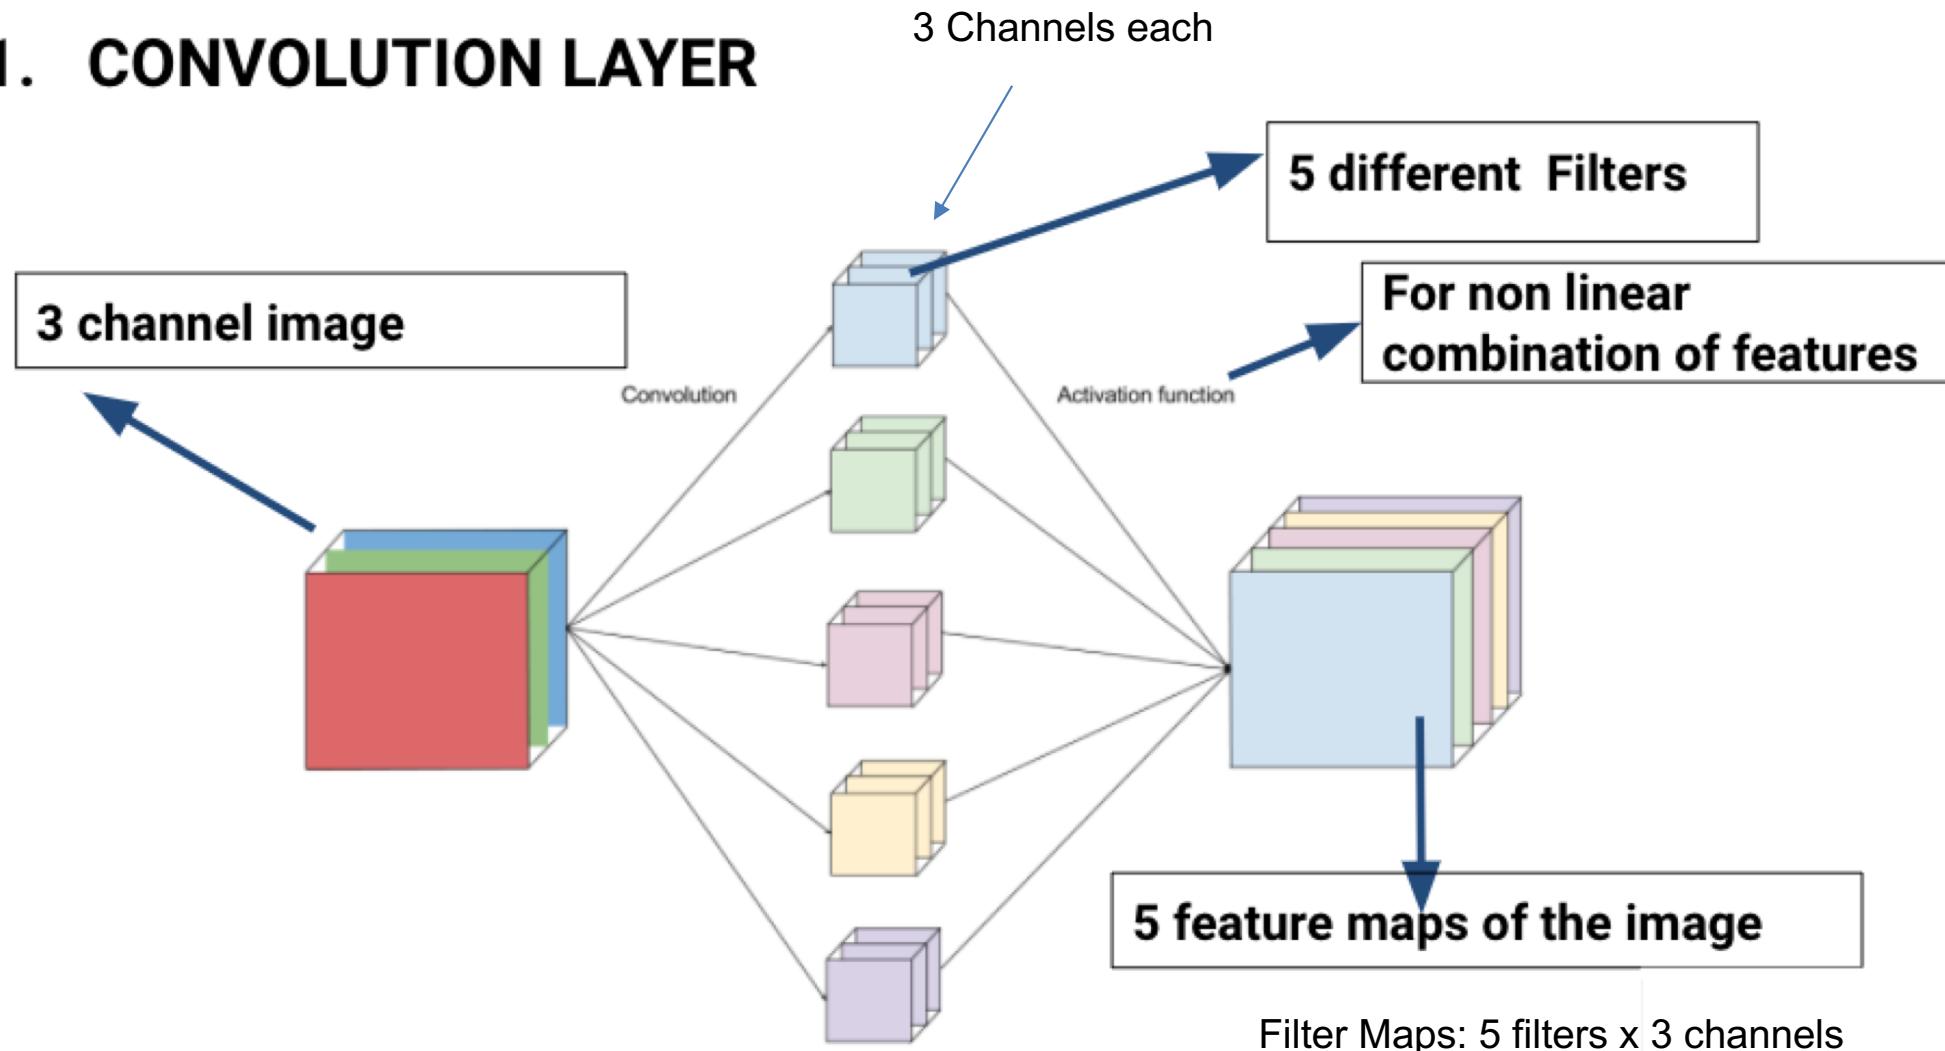


1. CONVOLUTION LAYER

For an image with **3 channels** and **multiple filters**

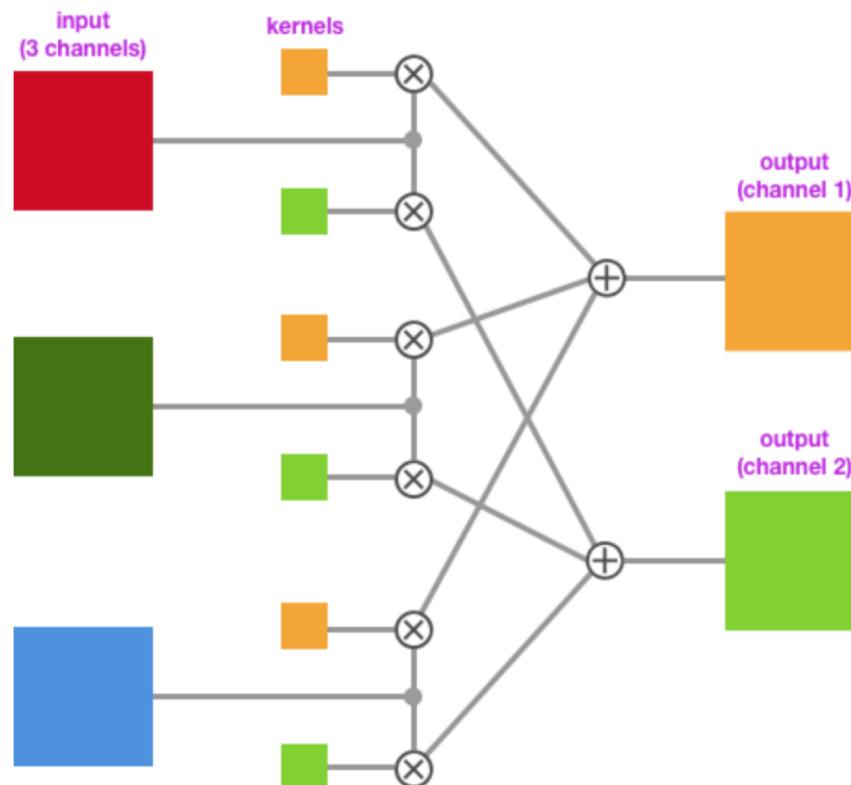


1. CONVOLUTION LAYER

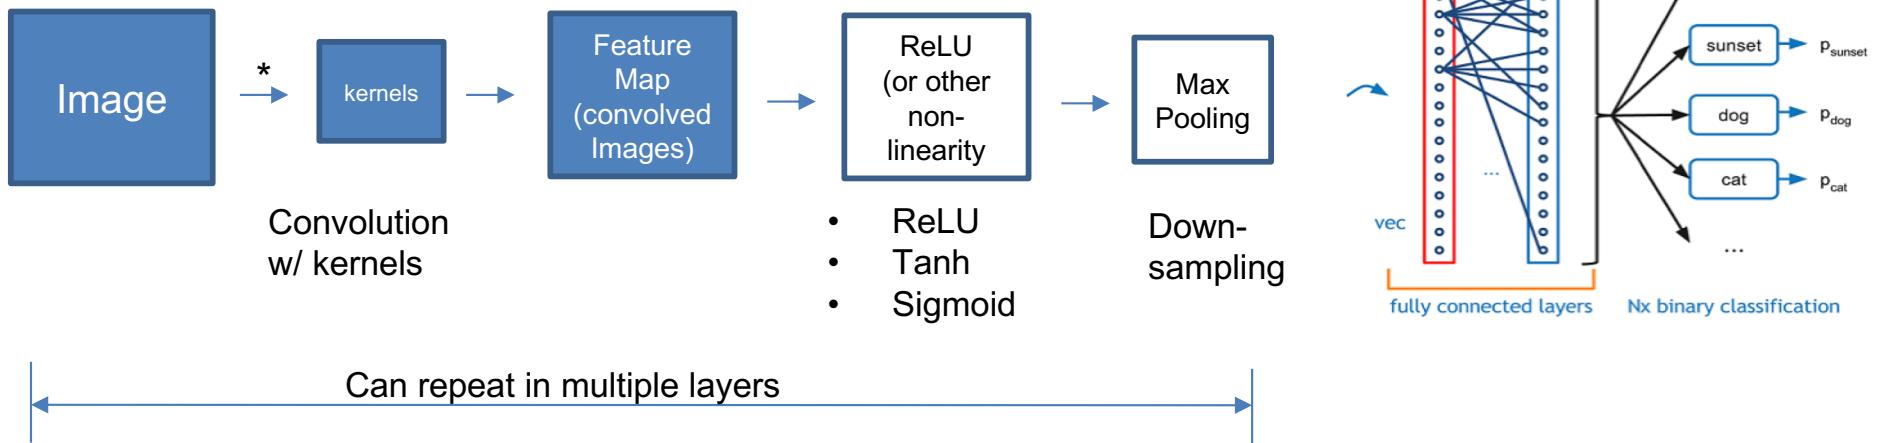


1. CONVOLUTION LAYER

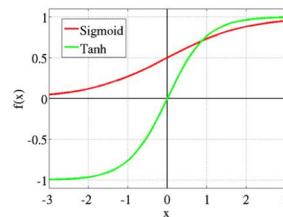
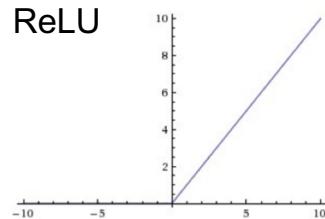
A detail: It's possible to sum the convolutions of 2 or more kernels



- 1) Convolution and Non-linearity Blocks First
- 2) Then Fully Connected Layers leading to prediction

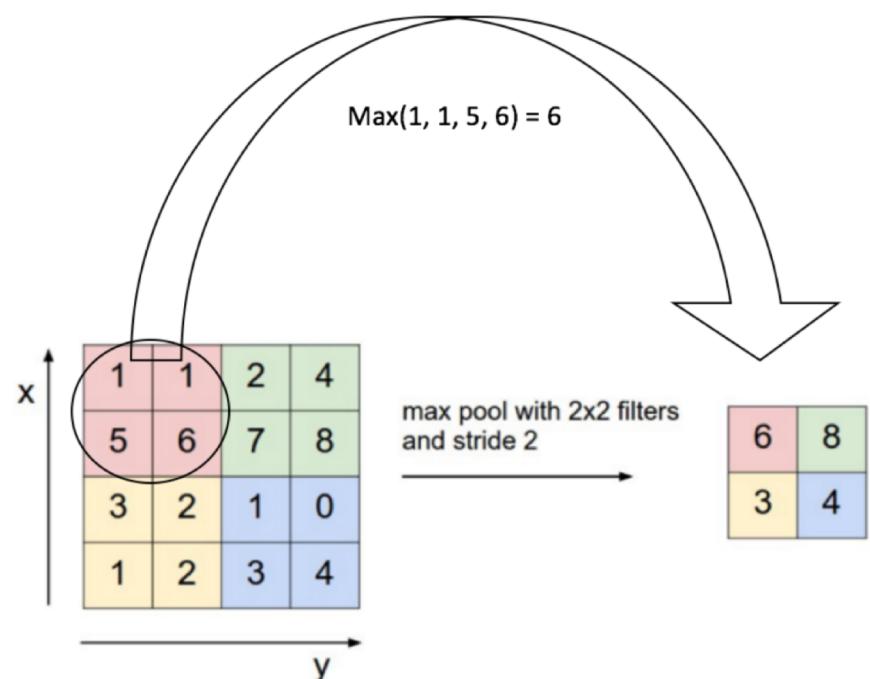


ReLU

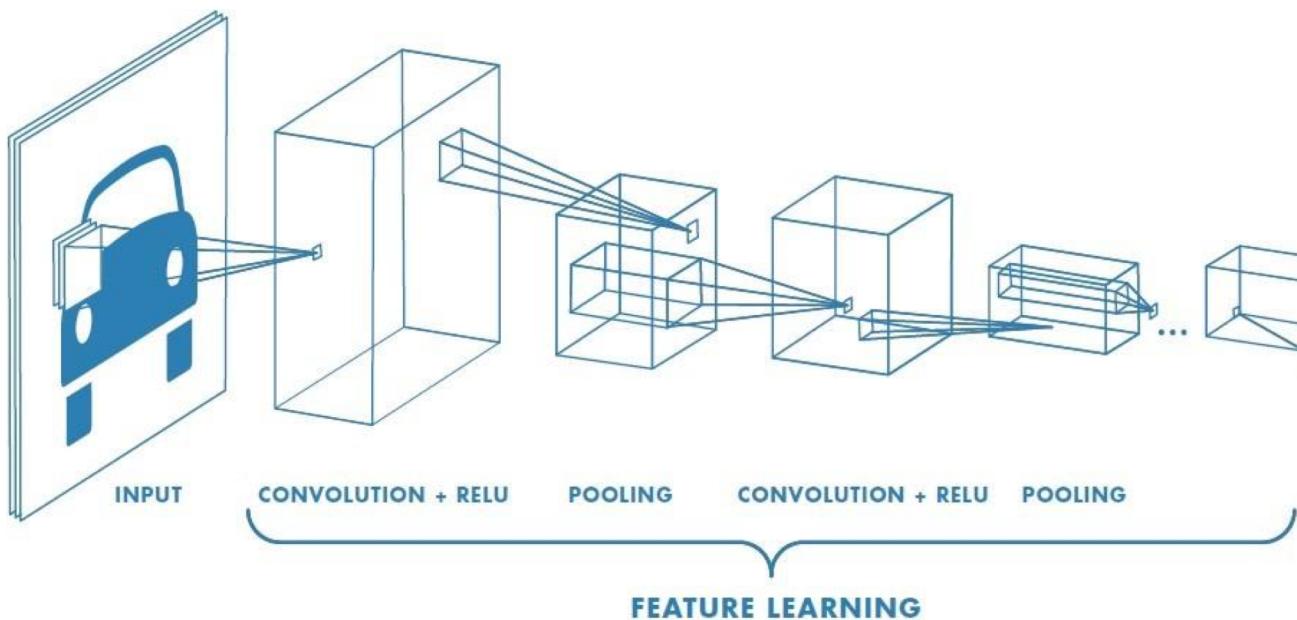


2. POOLING LAYER

1. Reduces the spatial dimensions (Width x Height) of the Input Volume
- 'down-sampling'
2. Usually max or average pooling is done.



Putting the CNN Together



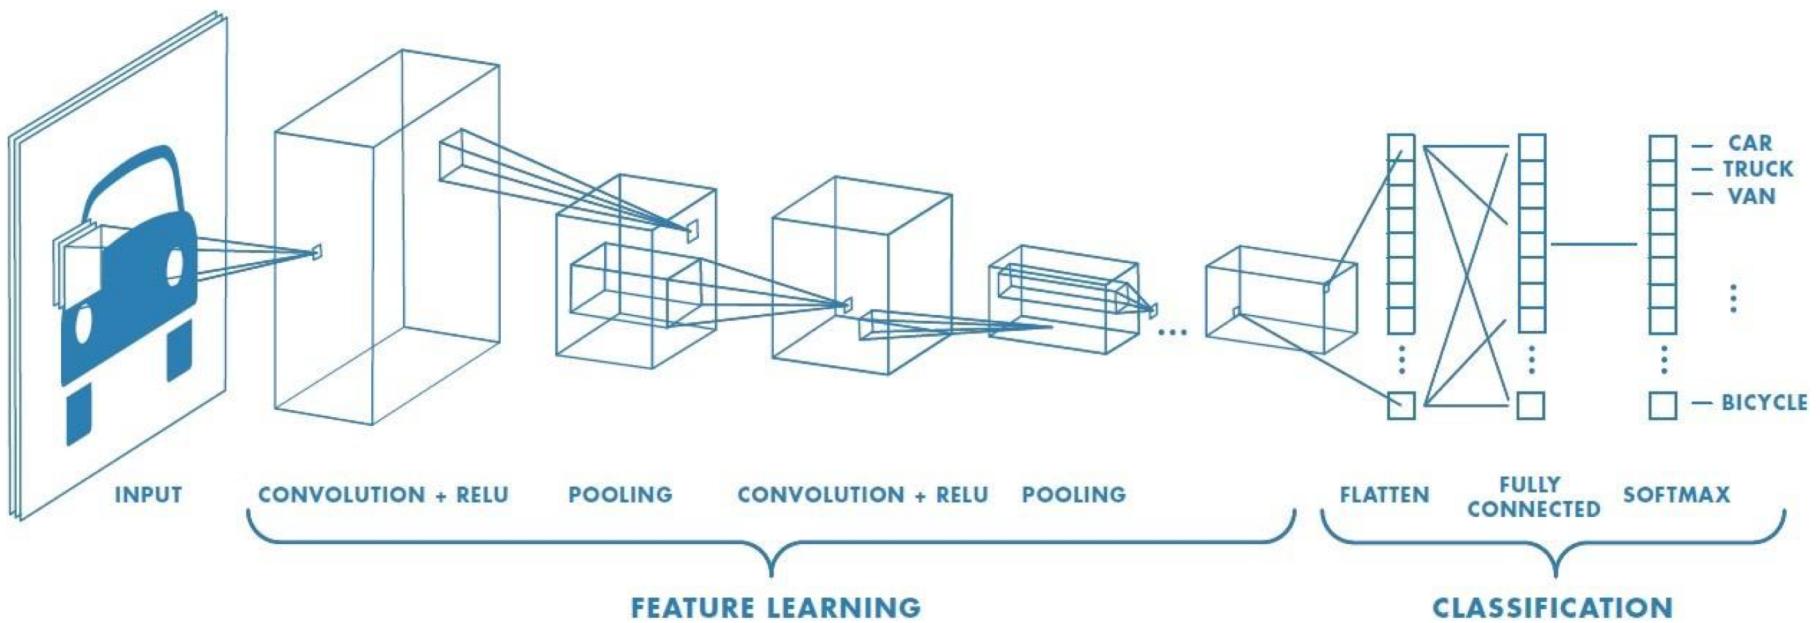
<https://www.mathworks.com/discovery/convolutional-neural-network.html>

Pooling:

- Reduces the spatial dimensions (Width x Height) of the Input Volume
 - 'down-sampling'
- It does not affect the depth dimension of the Volume.
- Usually placed after the Convolutional layer.
- The decrease in size leads to less computational overhead for the upcoming layers of the network

Data X

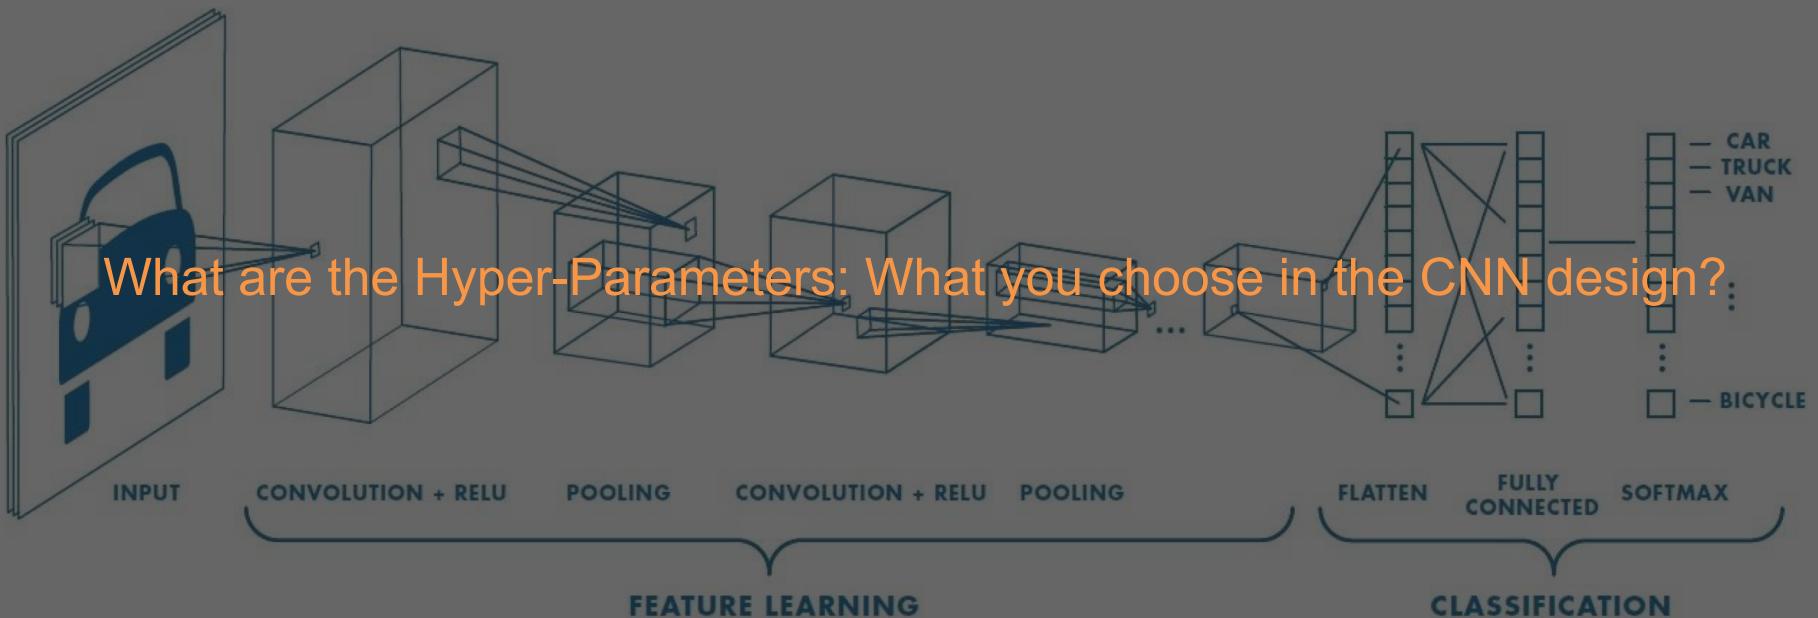
Putting the CNN Together with the Fully Connected Layer



<https://www.mathworks.com/discovery/convolutional-neural-network.html>



What are the Parameters—What the Algorithm Chooses? Putting the CNN together With the Fully Connected Layer



<https://www.mathworks.com/discovery/convolutional-neural-network.html>



Putting the CNN Together with the Fully Connected Layer What are the Parameters – What the Algorithm Chooses?

The filter or kernel itself!

What are the Hyper-Parameters: What you choose in the CNN design?

Number of kernels (features)
Sizes of the features ...
Number of Layers

Pooling Window Size, Pooling Stride

CLASSIFICATION

Fully Connected Layer: Number of Neurons, Number of Outputs

<https://www.mathworks.com/discovery/convolutional-neural-network.html>

Data X

End of Section

Data^X